Project Title: Build & Evaluate a Linear Regression Model — California Housing Dataset

Author Name: Venkata Durga Sai D

Tools Used: Python, pandas, scikit-learn, VS Code, Streamlit

Date: February 24, 2026

Project Description: This report details the end-to-end development of a Machine Learning pipeline designed to predict median housing prices across various neighborhoods in California. By leveraging historical census data, this project establishes a baseline predictive model, serving as a foundational step for more advanced real estate valuation systems.

Introduction

In the realm of Machine Learning, supervised learning is a paradigm where an algorithm learns from labeled training data to make predictions on unseen data. When the target variable we wish to predict is a continuous number—such as temperature, sales figures, or in this case, house prices—the problem is classified as a regression task.

The primary objective of this project is to build a reliable predictive model capable of estimating the median house value for a given block group in California. To achieve this, we employed Linear Regression, a fundamental statistical and machine learning technique. Linear Regression was selected because it provides a highly interpretable baseline; it assumes a linear relationship between the input variables and the single target variable, allowing us to clearly understand how each feature impacts the housing price.

Dataset Overview

The project utilizes the widely recognized California Housing Dataset, which was originally derived from the 1990 U.S. Census. It provides comprehensive demographic and geographical information aggregated at the block group level (a block group typically encompasses 600 to 3,000 people).

The dataset comprises several critical input features:

- **MedInc:** The median income of households within the block group.
- **HouseAge:** The median age of the houses in the block group.
- **AveRooms:** The average number of rooms per household.
- **AveBedrms:** The average number of bedrooms per household.
- **Population:** The total number of people residing in the block group.
- **AveOccup:** The average number of individuals living in a household.

- **Latitude & Longitude:** The geographical coordinates defining the location of the block group.

The target variable for our prediction is **MedHouseVal**, which represents the median house value for California districts, expressed in hundreds of thousands of dollars ($100,000s).

Exploratory Data Analysis (EDA)

Before modeling, an Exploratory Data Analysis (EDA) was conducted to uncover underlying patterns, spot anomalies, and understand feature relationships. Initial data profiling confirmed that the dataset was clean, with no missing or null values requiring imputation.

A correlation heatmap was generated to quantify the linear relationships between all variables. The most significant observation was the strong positive correlation between Median Income (MedInc) and the target variable (MedHouseVal). This intuitively aligns with real-world economics, where higher-income neighborhoods generally feature higher property values.

Furthermore, analyzing the feature distributions revealed that variables like Population and Average Occupancy exhibited right-skewed distributions, indicating the presence of localized dense populations or larger household sizes. The target variable itself showed a slight cap at the maximum value, an artifact of the census data collection process.

Model Training Methodology

To ensure our model could generalize well to new, unseen data, we adopted a standard evaluation strategy. The dataset was partitioned using a train-test split, randomly allocating 80% of the data to the training set and retaining the remaining 20% for testing.

Using the Scikit-Learn library, a Multiple Linear Regression model was instantiated and trained on the 80% subset. During the training phase, the algorithm utilized Ordinary Least Squares (OLS) optimization to calculate the optimal weights (coefficients) for each feature. This process systematically minimizes the residual sum of squares—the difference between the actual house prices in the training data and the prices predicted by the model's linear equation.

Evaluation Metrics & Results

To objectively assess the model's predictive performance on the 20% test set, we utilized three standard regression metrics:

- **Mean Absolute Error (MAE):** This metric represents the average absolute difference between the predicted values and the actual values. It provides a straightforward interpretation of the expected error margin in the target variable's native units.
- **Root Mean Squared Error (RMSE):** RMSE measures the standard deviation of the prediction errors. Because it squares the errors before averaging them, RMSE heavily penalizes larger discrepancies, making it a strict measure of model accuracy.

- **R-squared (R^2):** The coefficient of determination indicates the proportion of the variance in the target variable that is predictable from the independent variables. A score closer to 1.0 represents a perfect fit.

Model Performance Summary:

- MAE = [INSERT MAE]
- RMSE = [INSERT RMSE]
- R^2 Score = [INSERT R2]

Visual diagnostics were also performed. An "Actual vs. Predicted" scatter plot was plotted to observe the alignment of our predictions against a perfect-prediction reference line.

Additionally, a residual analysis plot was created by graphing the prediction errors against the predicted values. While the model captured the general trend, the residual plot indicated some heteroscedasticity (unequal variance of errors), suggesting that the linear assumption does not perfectly capture the complex dynamics of the California housing market.

Model Deployment

To transition the model from a theoretical script to a practical tool, it was serialized and saved to disk as a .pkl (pickle) file. This allows the trained algorithm to be loaded into memory without needing to be retrained.

Subsequently, a modern, interactive web application was developed using the Streamlit framework. This user interface allows end-users—such as real estate agents or prospective buyers—to input custom neighborhood characteristics (like median income or house age) via intuitive sliders and input boxes. The application then processes these inputs through the loaded pickle model to instantly return a formatted, estimated property value.

Conclusion & Future Improvements

This project successfully established a complete end-to-end Machine Learning pipeline. We ingested raw census data, performed exploratory analysis, trained a Linear Regression model, evaluated its performance, and deployed it via a Streamlit web interface. While Linear Regression served as an excellent, interpretable baseline, housing data is inherently complex and non-linear.

To improve predictive accuracy in future iterations, the following steps are recommended:

1. **Advanced Algorithms:** Transitioning to robust tree-based models like Random Forest or Gradient Boosting (XGBoost/LightGBM) to capture non-linear relationships and feature interactions.
2. **Regularization:** Implementing Ridge or Lasso regression to penalize overly complex models and handle potential multicollinearity between features.

3. **Feature Engineering & Scaling:** Applying transformations (like logarithmic scaling) to skewed features such as Population, and standardizing all features using a StandardScaler to improve convergence.
4. **Cloud Deployment:** Hosting the Streamlit application on cloud platforms such as AWS, Google Cloud, or Streamlit Community Cloud to make the tool accessible to a broader audience via a public URL.